

IMAGE SEGMENTATION BY SUPERPIXELS

T.Dumont^a, B.Figliuzzi^b

a. MINES ParisTech, theo.dumont@mines-paristech.fr

b. MINES ParisTech CMM, bruno.figliuzzi@mines-paristech.fr

Key-words:

deep learning; convolutional neural networks; image segmentation

Abstract:

In this paper, we study different options to improve the performance of a deep learning convolutional neural network

Contents

1	Introduction	3
1.1	Segmentation	3
1.2	Superpixels	3
1.3	Ce qu'est une bonne superpixelsegmentation	4
1.3.1	Metrics	4
1.3.2	Autres algorithmes	4
1.4	Motivations/ambitions	4
2	Dataset generation	4
2.1	COCO dataset	4
2.1.1	The COCO dataset	4
2.1.2	Characteristics	4
2.2	Eikonal	5
2.3	Notre utilisation de eikonal	5
3	The model	5
3.1	Approach	5
3.2	Network architecture	5
3.2.1	Layers definitions	5
3.2.2	Chen	7
3.2.3	UNet	8
3.2.4	Chen + UNet	8
3.3	Total Variation (TV) Loss	8
3.3.1	MSE	8
3.3.2	TV	8
3.4	Implementation	8
4	Expérience et résultats	9
4.1	Hyperparameters	9
4.1.1	Learning rate	9
4.1.2	Network size d	9
4.1.3	Number of epochs	9
4.1.4	TV regularization	9
4.1.5	Runs	9

4.2	Results on dataset	10
4.2.1	The dataset	10
4.2.2	Results	10
5	Conclusion/Discussion	10

Todo

-

1 Introduction

1.1 Segmentation

What it is `piche`



Figure 1: *An image and its segmented image*

Motivation

1.2 Superpixels

What it is `piche`



Figure 2: *A superpixel segmentation. From left to right: the original image, the original image with its calculated superpixels outlines and the resulting superpixel segmented image. Each pixel of a superpixel has only one color, the mean color of the original image over the superpixel region.*

Applications & Motivation `démarrer une segmentation`
fournir un support sur lequel faire de la classification (couleur/texture moyenne, etc)

1.3 Ce qu'est une bonne superpixelsegmentation

pour obtenir metriques des algos, `text`

1.3.1 Metrics

cf article <https://arxiv.org/pdf/1612.01601.pdf> let $S = S_{j=1}^K$ and $G = G_i$ be partitions of the same image $I : x_n \mapsto I(x_n)$, $1 \leq n \leq N$ S is a segmented image G is the ground truth

Boundary Recall - most commonly used metric to assess boundary adherence. - intersection entre les frontieres grossies et truth, donc pourcentage de vrais contours dectctes - Let $TP(G, S)$ be the number of true positive boundary pixels and $FN(G, S)$ be the number of false negative boudary pixels in the segmented image S .

$$\text{Rec}(G, S) = \frac{TP(G, S)}{TP(G, S) + FN(G, S)}$$

Undersegmentation Error

Compactness - evaluates the compactness of the superpixels.

$$\text{CO}(G, S) = \frac{1}{N} \sum_{S_j} |S_j| \frac{4\pi A(S_j)}{P(S_j)}$$

- the CO operator computes how close the area $A(S_j)$ of each superpixel S_j is from a circle with same perimeter $P(S_j)$.

1.3.2 Autres algorithmes

SLIC

1.4 Motivations/ambitions

Difficultés que l'on cherche à résoudre

Pas de vraie approche DL pour segmentation avec superpixels

Ambitions améliorer les métriques

2 Dataset generation

2.1 COCO dataset

2.1.1 The COCO dataset

COCO dataset¹, nb of images, examples

2.1.2 Characteristics

Piche

d'où la nécessité de faire des transformations sur la dataset pour uniformiser les tailles en entrée du réseau.

¹site de COCO

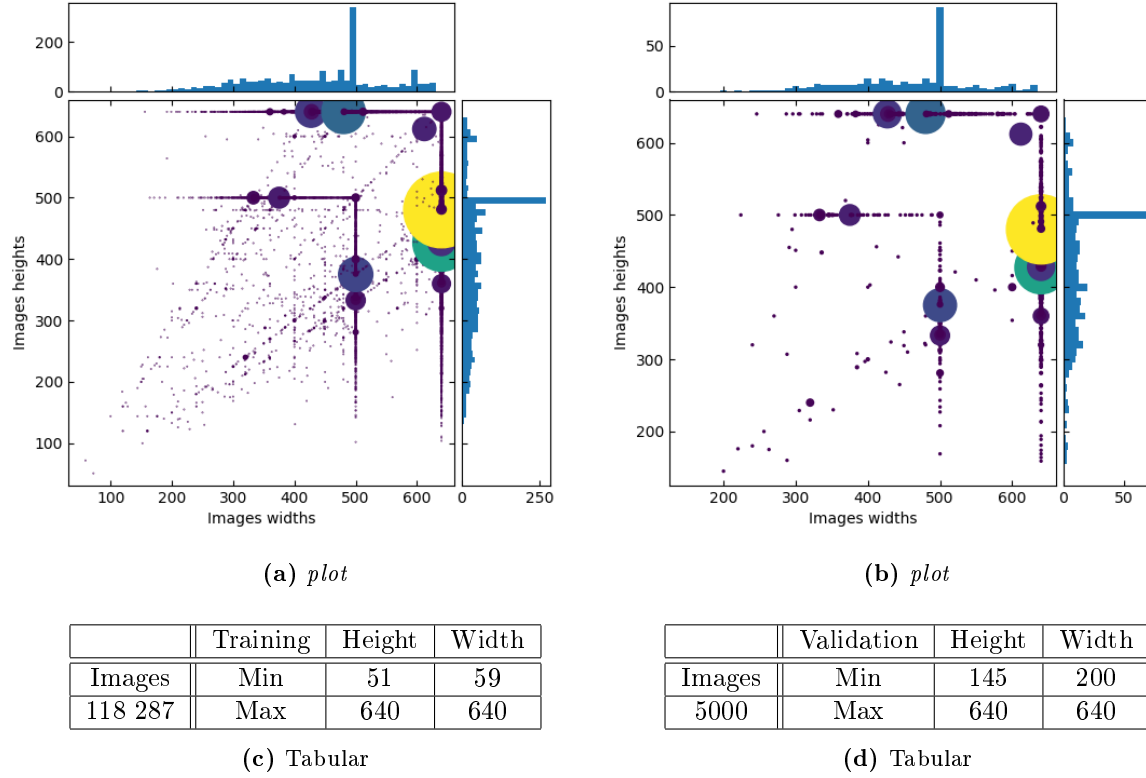


Figure 3: *Training and validation sets characterization*

2.2 Eikonal

2.3 Notre utilisation de eikonal

en plus réutilisé derrière sur image qui sort du réseau
faire un petit résumé

3 The model

3.1 Approach

description générale de l'approche (NN puis eikonal)

3.2 Network architecture

3.2.1 Layers definitions

Dilated convolution We consider a layer $L = (L_j)_{j \in \llbracket 1, w \rrbracket}$, w being the number of feature maps L_j of L . We also consider $K = (K_{i,j})_{i,j}$, each $K_{i,j}$ being a 3×3 convolutional kernel. The dilated convolution operation of $K_{i,j}$ on L_j is denoted by $L_j *_r K_{i,j}$, r being the dilation parameter.

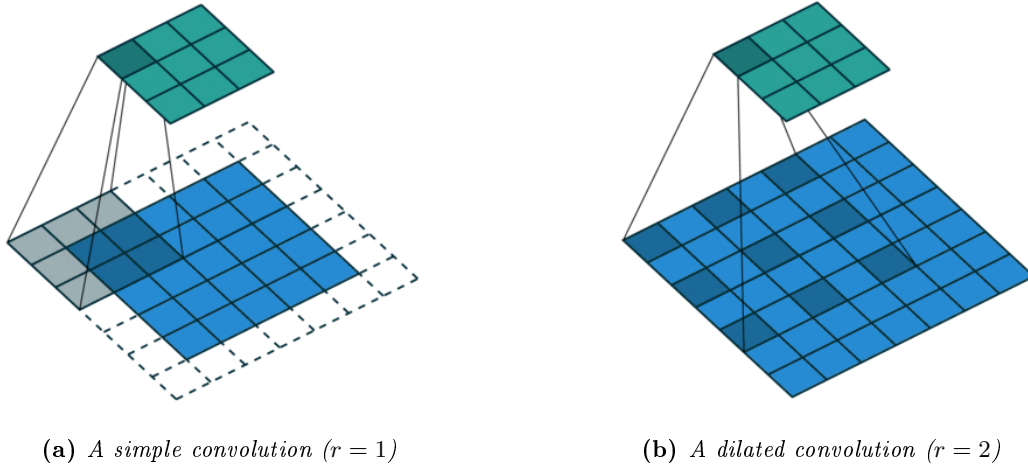


Figure 4: Illustration of two types of convolutions

The output $C(x)$ of a pixel x is:

$$\begin{aligned}
 C(x) &:= (L_j *_r K_{i,j})(x) \\
 &= \sum_{a+rb=x} L_j(a) K_{i,j}(b) \\
 &= \sum_b L_j(x - rb) K_{i,j}(b)
 \end{aligned}$$

and we recognize the simple convolution when $r = 1$.

A dilated convolution enables the network getting larger receptive fields while preserving the input resolution²

Adaptive Batch Normalization (ABN) As we have seen in (2.1.2), page 4, we need to normalize the data. We define the *adaptive normalization function* Ψ as:

$$\Psi(x) = a x + b BN(x),$$

where BN is the classic batch normalization³, defined as:

$$BN(x) = \frac{x - E[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta.$$

As such, Ψ combines identity mapping and batch normalization. a , b , γ and β are learned parameters⁴ by backpropagation. It allows the model to adapt to each dataset, choosing whether or not giving a big importance to the identity term and the normalization term.

Leaky rectifier (LReLU) In order to let our neural network model complex patterns in the data, we have to add a non-linear property to the model. It often is an activation function, such as a sigmoid or a tanh (Figure 5).

these activation functions are often used but they are bounded and their gradient is very low on the edges. Because we are going to manipulate high scalar values, we have to use an unbounded activation function, such as ReLU, $\Phi(x) = \max(0, x)$ (Figure 6a). But the issue with ReLU is that all the negative

²ref ?

³reference ?

⁴ref : https://pytorch.org/docs/stable/_modules/torch/nn/modules/batchnorm.html




Figure 5: *Illustration of two bounded rectifiers*

values become zero immediately, which decreases the ability of our model to train from the data. Hence the implementation of a *leaky rectifier*, LReLU (6b):

$$\Phi(x) = \max(\alpha x, x), \text{ with } 0 < \alpha < 1.$$




Figure 6: *Illustration of two unbounded rectifiers*

By implementing a Leaky Rectifier, we are able to take into account the negative valued pixels.

3.2.2 Chen

Context Aggregation Network (CAN) ⁵ blabla sur le RGB en entrée, RGB en sortie $I \rightarrow f(I)$

input I	\rightarrow	L^1	\rightarrow	\dots	\rightarrow	L^s	\rightarrow	\dots	\rightarrow	output (L^d)
$m \times n \times 3$		$m \times n \times w_1$				$m \times n \times w_s$				$m \times n \times 3$

Table 1: *Layers*

Architecture of a block Each block L_s is made of 3 layers:

1. A *dilated convolution*, $r_s = 2^s$

⁵reference

2. An adaptative batch normalization

3. A leaky rectifier (ReLU)

so that the content of an intermediate layer L^s can be computed from the content of the previous layer L^{s-1} :

$$L_i^s = \Phi \left(\Psi^s \left(b_i^s + \sum_j L_j^{s-1} *_{r_s} K_{i,j}^s \right) \right). \quad (1)$$

where ... is ...

and

$$L_j^{s-1} *_{r_s} K_{i,j}^s = \sum_{a+r_s b=x} L_j^{s-1}(a) K_{i,j}^s(b) \quad (2)$$

because of 3.2.1, page 5.

Layer	1	2	3	4	5	6	7
Convolution	3×3						
Dilation	1						
Batch Normalization	Yes						
LReLU	Yes						

Table 2: *Chen*

3.2.3 UNet

3.2.4 Chen + UNet

3.3 Total Variation (TV) Loss

3.3.1 MSE

$$L_{MSE} = \frac{1}{N} \sum_{i=1}^N |\hat{f}(I)_i - f(I)_i|^2$$

3.3.2 TV

6

Why In the process of creating an image segmentation neural network, we want the image output to have a lot of areas where the gradient is 0. We thus want to take into account this gradient $\nabla f(I)$ in the loss function [...]

Formula

$$L_{TV} = \frac{1}{N} \sum_{i=1}^N |\hat{f}(I)_i - f(I)_i|^2 + \frac{1}{N} \sum_{i=1}^N |(\nabla f(I))_i|^2$$

granting an improvement in the output image smoothness.

3.4 Implementation

The network was implemented with PyTorch⁷ and we used GPU acceleration [...] (pytorch), se renseigner (section assez courante) GPU acceleraation code sur github

⁶ref

⁷Repository can be found at <https://github.com/theodumont/superpixels-segmentation>.

4 Expérience et résultats

4.1 Hyperparameters

petit bilan des valeurs choisies évolution des paramètres a et b ?

4.1.1 Learning rate

lr_0	decay?	saturation?	d	TV?
0.001	No	No	7	No
0.01	No	No	7	No
0.01	$\times 0.5$ every 2 epochs	10^{-4}	7	No
0.001	$\times 0.5$ every 2 epochs	10^{-4}	7	No

Table 3: Runs for learning rate tuning




Figure 7: Tuning of learning rate

Constant value entraînement (lr, alpha) -> courbes de loss, et loss qui sature (cluster) d'où changement de lr au cours des epochs

Non constant value

4.1.2 Network size d

- Learning rate initialisé à 0.01, divisé par 2 toutes les 2 époques, saturation à $1e-4$ - Pas de régularisation TV CONCLUSION des run 3 à 5: Il est préférable de laisser $d=7$. Entre $d=6$ et $d=7$, l'amélioration semble relativement faible. bon intermédiaire entre temps de calcul et performances

4.1.3 Number of epochs

nb epochs: on le sélectionne en prenant le minimum de la validation loss

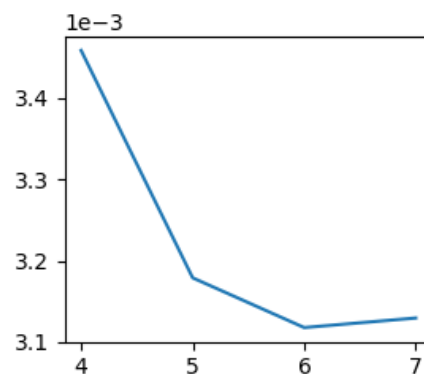
4.1.4 TV regularization

4.1.5 Runs

tableaux et graphes

d	4	5	6	7	8
$\text{loss} \times 10^3$	3.46	3.18	3.12	3.13	???

(a) Loss values on validation set



(b) Graph

Figure 8: Tuning of network size

4.2 Results on dataset

image originale → CNN → résultat du filtre dans eikonal → superpixels sans couleurs + couleur moyenne pour chaque spp de l'image originale cf results/images

4.2.1 The dataset

BSD dataset, BSD fait à la main et pas à l'arrache comme COCO

4.2.2 Results

metrics Here are the previously defined metrics of some well-known superpixel segmentation algorithms.

Algorithm		BR	UE	CO
image analysis	EI [8]			
	EIT [9]			
	WS [10]			
neural networks	SLIC [11]			
	Ours			

Table 4: Comparisons of metrics on the BSD dataset for different superpixel segmentation algorithms

We use the ?? algorithm as a reference to evaluate the performances of our model.

5 Conclusion/Discussion

On a présenté un nouveau...

On a prouvé...

Il reste à faire...

relire tous les mails pour avoir toutes les infos sur performances etc

Special thanks

Sources

- [1] [1] C. Smith, J.C. Green, Titre de l'article, Titre du journal, 10 (2009) 55-72
- [2] M. Truk, C. Bidul. Titre du bouquin, John Wiley and Sons, New York, 1973
- [3] P. Machin, Titre de la thèse, Thèse, Université Poitiers, 1992
- [4] D. Pierre, J.-P. Paul, B. Jacques, Titre communication, in: D. Editor, G. Editeur, (éd.), Proceedings of Conference XXX , Publisher, Paris, France, 1995, pp. 3–6